

# Log-Sum-Exp and Gumbel Intersection

## Motivation

**The “smooth maximum” puzzle:** What if you want to take the maximum of two numbers, but you need the result to be **smooth** (differentiable) for optimization? The hard maximum  $\max(x, y)$  has a sharp corner—it’s not differentiable at  $x = y$ .

Enter log-sum-exp:  $\beta \log(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}})$ . This is a smooth approximation to the maximum that becomes exact as  $\beta \rightarrow 0$ . It’s like replacing a sharp corner with a smooth curve that gets sharper as you zoom in.

When computing box intersections, we need to find the maximum of two Gumbel-distributed coordinates. The remarkable fact—and here’s the beautiful part—is that the location parameter of this maximum is given by the log-sum-exp function. It’s not just a convenient approximation; it’s the **exact** answer when working with Gumbel distributions.

**The “magic trick”:** The log-sum-exp function appears naturally in the Gumbel-Max trick: if we add Gumbel noise to deterministic values and take the maximum, the resulting distribution’s location parameter is the log-sum-exp of the original values. This provides both a theoretical foundation and a practical computational tool for box intersection operations. It’s as if Gumbel distributions were designed to produce this elegant result.

The temperature parameter  $\beta$  controls the “softness” of the maximum: as  $\beta \rightarrow 0$ , we recover the hard maximum (sharp corner); as  $\beta \rightarrow \infty$ , we approach the arithmetic mean (completely smooth). This flexibility allows us to interpolate between deterministic and probabilistic behavior—like a dial that controls how “hard” or “soft” our maximum operation is.

**The Gumbel-Max trick:** This result is the foundation of the Gumbel-Max trick, a technique widely used in machine learning for sampling from categorical distributions. If we have  $k$  categories with log-probabilities  $\log p_1, \dots, \log p_k$ , we can sample by adding independent Gumbel noise to each and taking the argmax. The resulting distribution is exactly the categorical distribution with probabilities  $p_1, \dots, p_k$ . This trick connects Gumbel distributions to discrete sampling and provides a differentiable relaxation (Gumbel-Softmax) for gradient-based optimization. In our context, we’re using the same mathematical structure: adding Gumbel noise and taking maxima naturally produces log-sum-exp in the location parameter.

**Historical context:** The Gumbel-Max trick has deep roots in statistics and social sciences. The connection between Gumbel distributions and log-sum-exp was known in social sciences before its application to machine learning. The trick is also related to the reparameterization trick in variational inference: by separating deterministic parameters from stochastic noise, we can compute gradients through the sampling process. The Gumbel-Softmax distribution (Jang et al., 2016; Maddison et al., 2016) extends this by replacing the non-differentiable argmax with a differentiable softmax, controlled by a temperature parameter. This temperature parameter allows interpolation between discrete (low temperature) and continuous (high temperature) behavior, enabling effective gradient-based optimization of discrete latent variables.

## Definition

The **log-sum-exp function** with temperature  $\beta > 0$  is:

$$\text{lse}_{\beta(x,y)} = \beta \log(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}})$$

We denote this as  $\text{lse}_{\beta(x,y)}$  to emphasize the temperature parameter  $\beta$ , which controls the “softness” of the operation. For Gumbel boxes, intersection coordinates are computed using log-sum-exp to determine the location parameters of the resulting Gumbel distributions.

## Statement

**Theorem (Gumbel-Max Property).** If  $G_1, G_2 \sim \text{Gumbel}(0, \beta)$  are independent **centered** Gumbel random variables (location parameter 0), then:

$$\max(x + G_1, y + G_2) \sim \text{Gumbel}\left(\text{lse}_{\beta(x,y)}, \beta\right)$$

The location parameter of the maximum is the log-sum-exp of the input locations  $x$  and  $y$ , while the scale parameter  $\beta$  remains unchanged. This is a manifestation of max-stability (see the Gumbel Max-Stability document): the maximum of Gumbel-distributed variables remains Gumbel-distributed, with the location parameter determined by log-sum-exp.

## Proof

The CDF of  $\max(x + G_1, y + G_2)$  is:

$$P(\max(x + G_1, y + G_2) \leq z) = P(x + G_1 \leq z \wedge y + G_2 \leq z)$$

Since  $G_1$  and  $G_2$  are independent:

$$= P(G_1 \leq z - x) * P(G_2 \leq z - y)$$

For  $\text{Gumbel}(0, \beta)$ , the CDF is  $F(z) = e^{-e^{-\frac{z}{\beta}}}$ , so:

$$= e^{-e^{-\frac{z-x}{\beta}}} * e^{-e^{-\frac{z-y}{\beta}}} = e^{-\left(e^{-\frac{z-x}{\beta}} + e^{-\frac{z-y}{\beta}}\right)}$$

Factoring out  $e^{-\frac{z}{\beta}}$  from the sum:

$$= e^{-e^{-\frac{z}{\beta}} \left(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}}\right)}$$

Now we manipulate the exponent to match the Gumbel CDF form. We want to write this as  $e^{-e^{-\frac{z-\mu}{\beta}}}$  for some location parameter  $\mu$ . To do this, we factor:

$$e^{-\frac{z}{\beta}} \left(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}}\right) = e^{-\frac{z - \beta \ln(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}})}{\beta}}$$

Therefore:

$$= e^{-e^{-\frac{z - \beta \ln(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}})}{\beta}}}$$

This is the CDF of  $\text{Gumbel}\left(\beta \ln\left(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}}\right), \beta\right)$ . Since  $\beta \ln\left(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}}\right) = \text{lse}_{\beta(x,y)}$ , we have:

$$\max(x + G_1, y + G_2) \sim \text{Gumbel}\left(\text{lse}_{\beta(x,y)}, \beta\right)$$

**Why log-sum-exp appears:** This result emerges naturally from the structure of Gumbel distributions. The Gumbel CDF has the form  $e^{-e^{-\frac{x-\mu}{\beta}}}$ , and when we multiply two such CDFs (for the maximum), the exponential structure leads to a sum of exponentials in the exponent. Taking the logarithm of this sum (after appropriate scaling) gives the log-sum-exp function. This is why log-

sum-exp is the “natural” operation for Gumbel distributions, just as addition is natural for normal distributions.

The appearance of log-sum-exp is not arbitrary—it’s a consequence of the exponential family structure of Gumbel distributions. When we compute the CDF of the maximum  $P(\max(x + G_1, y + G_2) \leq z)$ , we get  $P(x + G_1 \leq z) * P(y + G_2 \leq z) = e^{-e^{-\frac{z-x}{\beta}}} * e^{-e^{-\frac{z-y}{\beta}}}$ . The product of exponentials becomes a sum in the exponent:  $e^{-\left(e^{-\frac{z-x}{\beta}} + e^{-\frac{z-y}{\beta}}\right)}$ . To match the Gumbel CDF form  $e^{-e^{-\frac{z-\mu}{\beta}}}$ , we need  $e^{-\frac{z-\mu}{\beta}} = e^{-\frac{z-x}{\beta}} + e^{-\frac{z-y}{\beta}}$ . Factoring out  $e^{-\frac{z}{\beta}}$  and solving for  $\mu$  gives  $\mu = \beta \ln(e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}})$ , which is exactly the log-sum-exp function. This algebraic manipulation reveals why log-sum-exp is the “natural” operation for Gumbel distributions.

## Numerical Stability

Direct computation of  $e^{\frac{x}{\beta}} + e^{\frac{y}{\beta}}$  can overflow when  $\frac{x}{\beta}$  or  $\frac{y}{\beta}$  are large. Similarly, it can underflow when both are very negative, leading to loss of precision. The stable form is:

$$\text{lse}_{\beta(x,y)} = \max(x, y) + \beta \log\left(1 + e^{-|x-y|/\beta}\right)$$

This is the log-sum-exp trick, a fundamental technique in numerical computing. By shifting the computation by the maximum value, we ensure that the largest exponentiated term is  $e^0 = 1$ , preventing overflow. The correction term is bounded:  $0 \leq \beta \log\left(1 + e^{-|x-y|/\beta}\right) \leq \beta \log 2$ , ensuring numerical stability.

**Why this works:** The key insight is that we can shift all values by an arbitrary constant  $c$  without changing the result:  $\log\left(\sum_i e^{x_i}\right) = c + \log\left(\sum_i e^{x_i - c}\right)$ . By choosing  $c = \max_i x_i$ , we ensure that at least one term in the sum is  $e^0 = 1$ , and all other terms are  $\leq 1$ , preventing overflow. This trick is essential when working with log-probabilities in statistical modeling, where values can have arbitrary scale depending on the likelihood function and number of data points. The same principle applies to our temperature-scaled version: by factoring out  $e^{\frac{\max(x,y)}{\beta}}$ , we work with bounded terms that cannot overflow.

## Application to Box Intersection

For Gumbel boxes, intersection coordinates are computed using log-sum-exp to determine the location parameters:

- **Intersection minimum:**  $z_{\{\text{cap}, i\}} \sim \text{MaxGumbel}\left(\text{lse}_{\beta(\mu_{z,i}^A, \mu_{z,i}^B)}, \beta\right)$  This is the maximum of the two boxes’ minimum coordinates. The log-sum-exp gives the location parameter of the resulting MaxGumbel distribution.
- **Intersection maximum:**  $Z_{\{\text{cap}, i\}} \sim \text{MinGumbel}\left(\text{lse}_{\beta(\mu_{Z,i}^A, \mu_{Z,i}^B)}, \beta\right)$  This is the minimum of the two boxes’ maximum coordinates. Again, log-sum-exp provides the location parameter.

This construction preserves the Gumbel distribution family through intersection operations, maintaining algebraic closure (see the Gumbel Max-Stability document). The scale parameter  $\beta$  remains unchanged, ensuring consistency with the original box definitions.

## Limits

- As  $\beta \rightarrow 0$ :  $\text{lse}_{\beta(x,y)} \rightarrow \max(x, y)$  (hard maximum). The correction term  $\beta \log\left(1 + e^{-|x-y|/\beta}\right) \rightarrow 0$ , recovering the deterministic maximum. This limit is crucial: it shows that log-sum-exp is a smooth approximation to the maximum function, making it useful for optimization problems where we need differentiability but want behavior close to the hard maximum.

- As  $\beta \rightarrow \infty$ :  $\text{lse}_{\beta(x,y)} \rightarrow \frac{x+y}{2}$  (arithmetic mean). For large  $\beta$ , the exponential terms become similar, and the log-sum-exp approaches the average. This limit shows that log-sum-exp interpolates between the maximum (most selective) and the mean (most inclusive) operations.

The temperature parameter  $\beta$  controls the “softness” of the maximum operation, interpolating between deterministic (small  $\beta$ ) and probabilistic (large  $\beta$ ) behavior. This interpolation property is why log-sum-exp is sometimes called a “smooth maximum” or “soft maximum” in the machine learning literature. The function satisfies the bounds:  $\max(x, y) < \text{lse}_{\beta(x,y)} \leq \max(x, y) + \beta \log 2$ , showing that it’s always slightly larger than the hard maximum but never more than  $\beta \log 2$  away from it.

## Example

For  $x = 100.0$ ,  $y = 100.5$ ,  $\beta = 0.1$ :

**Unstable computation:**

$$\text{lse}_{\beta(x,y)} = 0.1 * \log(e^{1000} + e^{1005})$$

(overflows!)

**Stable computation:**

$$\text{lse}_{\beta(x,y)} = 100.5 + 0.1 * \log(1 + e^{-5}) \approx 100.5 + 0.0067 = 100.5067$$

The stable form avoids overflow by working with bounded correction terms.